

# Streaming SSM (Mamba-like) vs Kalman for 2D Tracking

A critical toy study on robustness under missing measurements

Code and materials: [https://github.com/segalshai/SSM-Style-Neural-Filters-for-2D-Tracking-A-Critical-Toy-Study/blob/main/Fast\\_Survey\\_Mamba\\_SSMS\\_for\\_Target\\_Tracking\\_with\\_Toy\\_Study.pdf](https://github.com/segalshai/SSM-Style-Neural-Filters-for-2D-Tracking-A-Critical-Toy-Study/blob/main/Fast_Survey_Mamba_SSMS_for_Target_Tracking_with_Toy_Study.pdf)

## Objective

Evaluate whether a small, causal state-space-model (SSM) style neural network ("Mamba-like") can act as a real-time tracker and provide any practical advantage over a classical constant-velocity (CV) Kalman filter. The focus is robustness when measurements are missing.

## Introduction and brief survey

Tracking is sequential state estimation under noisy and sometimes intermittent measurements. Classical Kalman filtering is attractive for stability and interpretability, but performance depends on the adequacy of the assumed motion prior (commonly constant velocity) and noise model. Under hard conditions—abrupt maneuvers, long measurement gaps, outliers, or non-stationary noise—a CV Kalman predictor can accumulate systematic error and diverge during occlusions because the prior becomes wrong and measurements are absent to correct it.

A classical response is to expand or switch motion models. The interacting multiple model (IMM) filter formalizes maneuvering as switching dynamics and remains a baseline when a small bank of models (e.g., CV/CT/CA) can plausibly cover target behavior [1]. However, performance still depends on the model set and tuning, and can degrade when maneuvers deviate from the assumed family.

Learning-augmented filters retain the predict/update structure while learning mismatch-sensitive components from data. KalmanNet preserves the Kalman flow and uses a compact recurrent module to learn gain-like updates that can improve estimation when dynamics and noise are only partially known [2,3]. The limitation is that learned behavior inherits the training distribution and any simulation-to-reality gap in sensing and dynamics [2,3].

Selective state-space sequence models (SSMs) motivate a related approach: learned, causal temporal models with efficient streaming inference. Mamba introduces a selective SSM architecture with linear-time scaling and a hardware-aware recurrent execution path, making it a candidate for high-rate streaming inference under strict compute/latency budgets [4]. In tracking, Mamba-like modules appear primarily as efficient temporal backbones (e.g., TrackingMamba) or as learned motion predictors inside tracking pipelines (e.g., MambaMOT) [5,6]. These works motivate SSMs mainly as learned motion priors rather than probabilistic filter replacements.

The remainder of this document reports a controlled toy study testing whether an SSM-inspired streaming block can learn turn-like motion priors and improve robustness specifically during measurement dropouts, relative to a baseline CV Kalman filter.

Date: January 12, 2026

## 1. Goal and research question

This project was a hands-on exploration of whether state-space-model (SSM) inspired neural sequence models ("Mamba-like" streaming blocks) have practical potential for tracking problems. The focus was not to "beat Kalman" in general, but to test a more specific hypothesis: **can an SSM-style model learn motion priors from data (e.g., line and turning segments) and then generalize to new trajectories composed from those primitives, especially when measurements are unreliable or missing?**

## 2. Experimental setup

**State:**  $x = [p_x, p_y, v_x, v_y]$  (4D). **Measurement:**  $z = [p_x, p_y]$  (2D). **Time step:**  $dt = 0.1$ . All experiments are synthetic and 2D. Motion is generated by concatenating primitives: straight segments (constant heading) and constant-turn arcs (left/right) with configurable speed and turn rate.

**Training data:** random sequences of primitives (typically 4-5 segments), with random  $v$  and  $\omega$  per sequence. **Generalization test:** evaluate on a held-out primitive pattern (unseen exact sequence) composed from the same primitive set. Measurement corruption modes were added progressively: Gaussian noise, outliers, sparse sampling, and forced block dropouts.

Figure 1 shows the typical "train on primitives, test on held-out combination" baseline (Gaussian noise only).

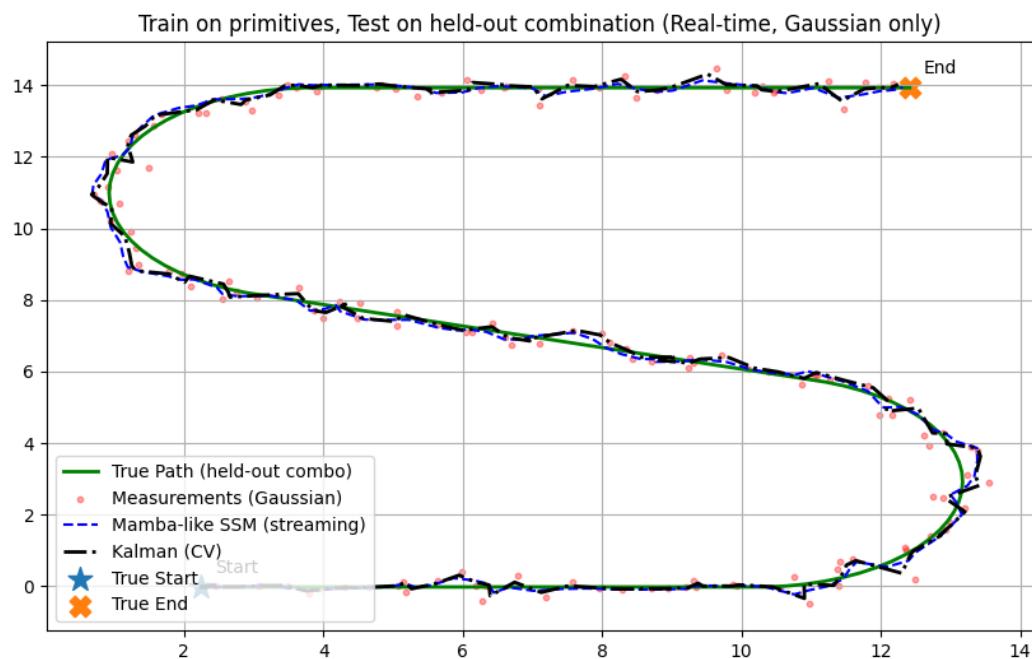


Figure 1. Train on motion primitives (line + arc), test on an unseen combination (Gaussian only).

### 3. Filters compared

#### 3.1 Kalman baseline

A standard **linear Kalman filter** with a **constant-velocity (CV)** motion model. Prediction runs every frame; measurement update is applied only when a measurement exists. Process noise Q and measurement noise R are scalar-scaled identity matrices (simple tuning). This baseline is intentionally minimal; it is not IMM/CT/UKF.

#### 3.2 Learned SSM-style streaming filter

A **streaming neural filter** built from two recurrent "Mamba-like" layers (simple gated state update) and trained end-to-end by backpropagation through time. To stay close to Kalman structure, the learned filter was implemented as:

**Predict:** start from current state  $x_t$ , apply CV propagation, and add a **learned residual dynamics** term (e.g., delta-v or a turn-rate proxy) that can run even when measurements are missing.

**Update:** when a measurement exists, compute innovation ( $z - \text{predicted position}$ ) and apply a **learned correction** to the full state (similar role to Kalman gain, but learned and nonlinear).

*Note:* this is SSM-inspired (streaming state update), but it is not a faithful implementation of the full Mamba architecture with explicit A,B,C matrices. The intent was educational and pragmatic: keep real-time step() inference and test whether learned state dynamics + learned measurement update can help in tracking regimes where CV is mismatched.

## 4. Training objective and optimization

Training is supervised: run the filter over a sequence and minimize position error to the known ground truth. The main loss was a **weighted SmoothL1** on position:

$L_{pos} = \text{mean}_t(w_t * \text{SmoothL1}(p_{hat\_t} - p_t))$ , where  $w_t$  is increased during missing-measurement frames to force the model to learn to "carry" motion through gaps.

Additional small regularizers were used in some runs: (1) a velocity consistency term from finite differences, and (2) (when using a turn-rate proxy) a smoothness penalty encouraging stable omega during gaps. Optimization used Adam (typical lr ~ 3e-3) for ~800-1200 epochs on small CPU batches.

**Important practical finding:** the learned filter can become numerically unstable (NaNs) when the scenario is made too hard (large gaps + strong turns + high learning rate). Clamping/tanh on residual outputs and modest regularization were needed.

## 5. Key experiments and results

Scenario	Meas. density	Forced gap	RMSE total (Learned / KF)	RMSE gap (Learned / KF)	MaxErr gap (Learned / KF)
Gaussian only (primitives -> held-out <del>calib</del> )	dense	none	0.182 / 0.208	n/a	n/a
Test-only outliers (no retraining)	dense	none	1.091 / 1.002	n/a	n/a
Sparse + long gap inside turn (K=5)	30 / 180	inside turn	1.887 / 3.524	4.181 / 5.584	6.030 / 10.171
Sparse + long gap inside turn (K=2)	76 / 180	inside turn	0.767 / 3.090	1.781 / 6.149	2.259 / 10.584

Interpretation: in the "Gaussian only" regime the learned filter is competitive but does not clearly dominate a well-tuned Kalman CV. In "test-only outliers" (no robustness training), both degrade similarly. The clearest advantage appears in the **forced long measurement gap occurring during a turning maneuver**: the learned model can maintain a plausible turn prior, while Kalman CV extrapolates the wrong motion and diverges.

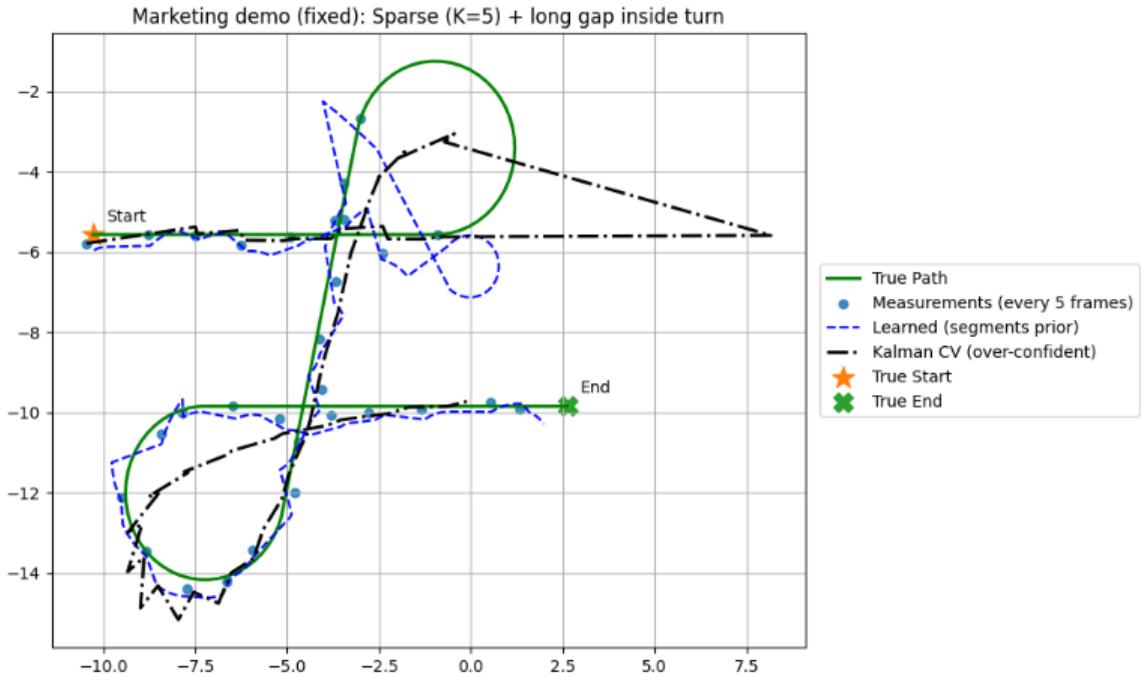


Figure 2. Sparse sampling ( $K=5$ ) with a forced block dropout inside a turn.

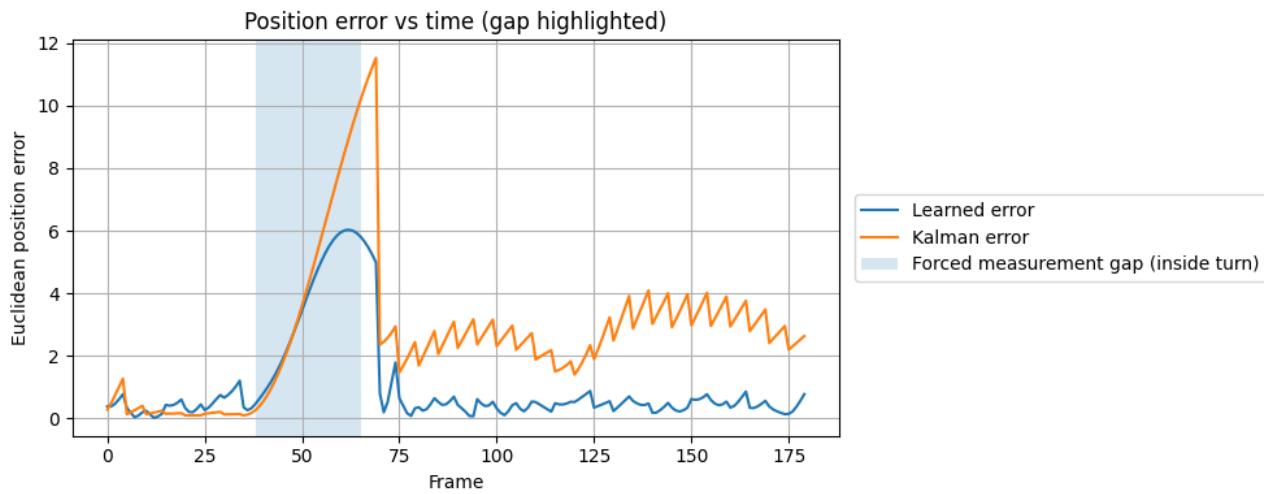


Figure 3. Position error vs time for Figure 2 (gap highlighted).

## 6. Focus result: K=2 (dense overall, but one long gap inside the turn)

With K=2 the tracker receives measurements frequently outside the forced gap. This makes the experiment visually clean: both methods look reasonable most of the time, and the gap inside a maneuver isolates the effect of the motion prior.

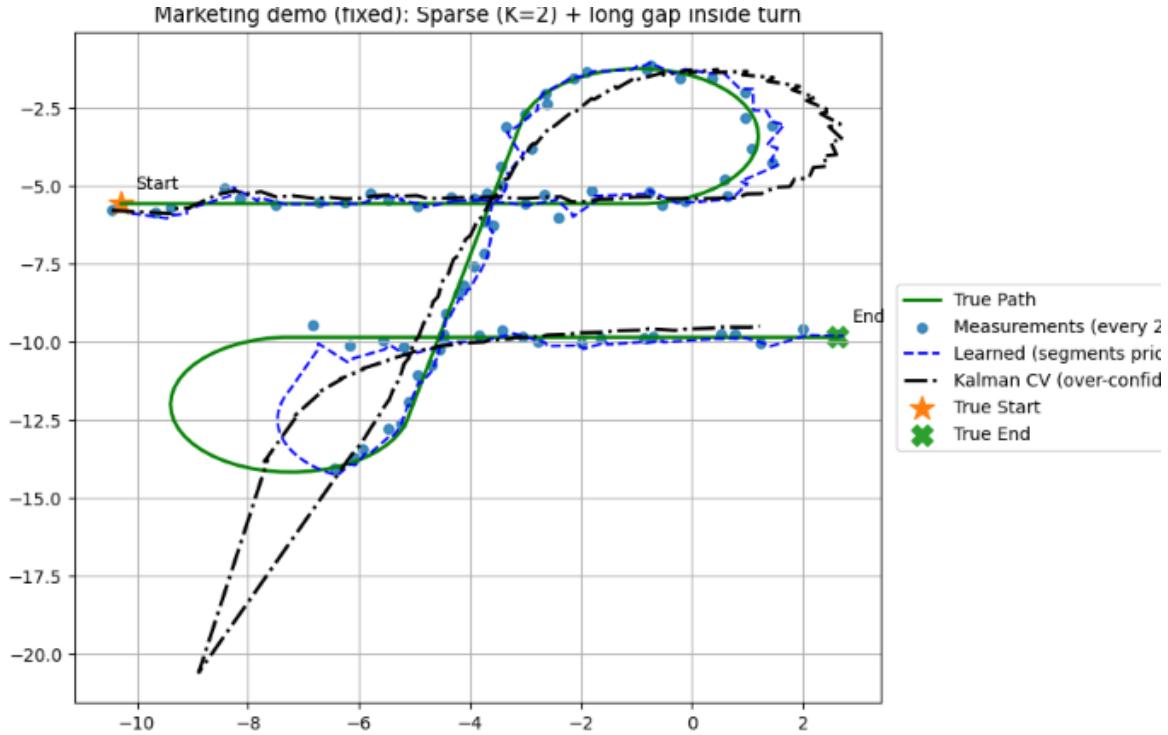


Figure 4. K=2 overall measurement rate, but a long forced gap inside the first turn.

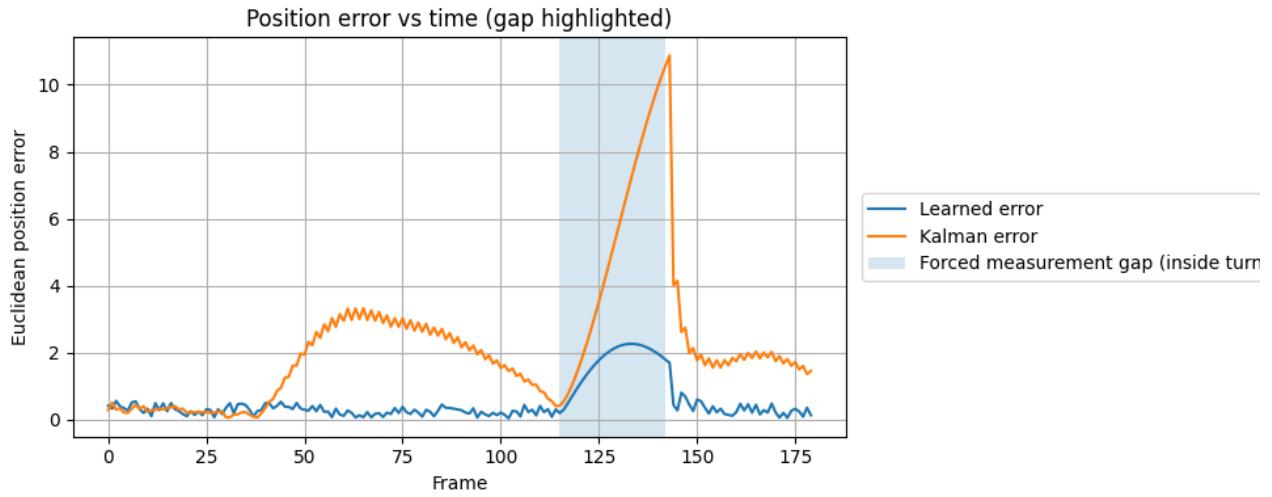


Figure 5. Error vs time for Figure 4: learned filter rises but stays bounded; Kalman CV spikes strongly.

## 7. Conclusions (critical and realistic)

**Do SSM-style networks have potential for tracking?** Yes, but the potential is situational. In this toy setting, the strongest evidence is that a learned streaming model can encode a **motion prior** (e.g., turning behavior) that helps during measurement gaps where a simplistic CV model is wrong. This

aligns with recent MOT literature that replaces/augments Kalman motion prediction with learned SSM modules.

**What did the learned model actually learn here?** Primarily a nonlinear measurement update and a residual dynamics term (turn tendency) conditioned on recent history. It did not discover new physics; it learned a data-driven maneuver model from the primitive distribution.

**Limitations:** (1) The learned model is not a full Mamba implementation with explicit A,B,C matrices; it is an SSM-inspired recurrent block. (2) Results depend heavily on training distribution and hyperparameters; instability (NaNs) can occur without output clamping. (3) The Kalman baseline is intentionally simple; stronger Kalman variants (CT/IMM/adaptive Q) would likely close much of the gap. (4) This is a synthetic single-target 2D toy problem; conclusions do not directly transfer to real sensors without additional work.

**Practical takeaway:** SSM-style learned filters are best viewed as **learned motion models** that can plug into a classical pipeline (or operate alongside it), especially for regimes with model mismatch and intermittent observations. For well-modeled regimes with frequent measurements, a tuned Kalman filter remains hard to beat on simplicity, stability, and interpretability.

Reference (motivation only): arXiv:2403.10826v2 (MambaMOT) discusses using an SSM module as a motion predictor in multi-object tracking.

## References

- [1] H. A. P. Blom and Y. Bar-Shalom, "The Interacting Multiple Model Algorithm for Systems with Markovian Switching Coefficients," *IEEE Transactions on Automatic Control*, 33(8), 1988. DOI: 10.1109/9.1299. <https://doi.org/10.1109/9.1299>
- [2] G. Revach, N. Shlezinger, R. J. G. van Sloun, and Y. C. Eldar, "KalmanNet: Data-Driven Kalman Filtering," *ICASSP 2021*. DOI: 10.1109/ICASSP39728.2021.9413750. <https://doi.org/10.1109/ICASSP39728.2021.9413750>
- [3] G. Revach, N. Shlezinger, R. J. G. van Sloun, and Y. C. Eldar, "KalmanNet: Neural Network Aided Kalman Filtering for Partially Known Dynamics," *arXiv:2107.10043*, 2021. <https://arxiv.org/abs/2107.10043>
- [4] A. Gu and T. Dao, "Mamba: Linear-Time Sequence Modeling with Selective State Spaces," *arXiv:2312.00752*, 2023. <https://arxiv.org/abs/2312.00752>
- [5] H.-W. Huang et al., "MambaMOT: State-Space Model as Motion Predictor for Multi-Object Tracking," *arXiv:2403.10826*, 2024. <https://arxiv.org/abs/2403.10826>
- [6] Q. Wang et al., "TrackingMamba: Visual State Space Model for Object Tracking," *IEEE JSTARS*, 2024 (see publisher record).

Code and materials: [https://github.com/segalshai/SSM-Style-Neural-Filters-for-2D-Tracking-A-Critical-Toy-Study/blob/main/Fast\\_Survey\\_Mamba\\_SSMS\\_for\\_Target\\_Tracking\\_with\\_Toy\\_Study.pdf](https://github.com/segalshai/SSM-Style-Neural-Filters-for-2D-Tracking-A-Critical-Toy-Study/blob/main/Fast_Survey_Mamba_SSMS_for_Target_Tracking_with_Toy_Study.pdf)